Probability - Lecture notes - Unit 01

Events and outcomes

01 Theory

Events and outcomes – informally

- An **event** is a *description* of something that can happen.
- An **outcome** is a *complete description* of something that can happen.

All outcomes are events. An event is usually a *partial* description. Outcomes are events given with a *complete* description.

Here 'complete' and 'partial' are within the context of the **probability model**.

- ⚠ It can be misleading to say that an 'outcome' is an 'observation'.
  - 'Observations' occur in the *real world*, while 'outcomes' occur in the *model*.
  - To the extent the model is a good one, and the observation conveys *complete* information, we can say 'outcome' for the observation.

Notice:

- Because outcomes are *complete*, no two distinct outcomes could *actually happen* in a run of the experiment being modeled.

When an event happens, the *fact* that it has happened constitutes **information**.

Events and outcomes – mathematically

- The **sample space** is the *set of possible outcomes*, so it is the set of the complete descriptions of everything that can happen.
- An **event** is a *subset* of the sample space, so it is a *collection of outcomes*.

- For mathematicians: some "wild" subsets are not *valid* events. Problems with infinity and the continuum...

Notation

- Write S for the set of possible outcomes, $s \in S$ for a single outcome in S.
- Write $A, B, C, \dots \subset S$ or $A_1, A_2, A_3, \dots \subset S$ for some events, subsets of S.
- Write $\mathcal{F}$ for the collection of all events. This is frequently a *huge* set!
- Write $|A|$ for the **cardinality** or *size* of a set A, i.e. the *number of elements it contains*.

Using this notation, we can consider an *outcome itself as an event* by considering the "singleton" subset $\{\omega\} \subset S$ which contains that outcome alone.

02 Illustration

Example - Coin flipping

Flip a fair coin two times and record both results.

- *Outcomes:* sequences, like HH or TH.
- *Sample space:* all possible sequences, i.e. the set $S = \{HH, HT, TH, TT\}$.
- *Events:* for example:
  - $A = \{HH, HT\} =$ "first was heads"
  - $B = \{HT, TH\} =$ "exactly one heads"
  - $C = \{HT, TH, HH\} =$ "at least one heads"

With this setup, we may combine events in various ways to generate other events:

- *Complex events:* for example:
  - $A \cap B = \{HT\}$, or in words:

    $$\text{"first was heads" AND "exactly one heads"} = \text{"heads-then-tails"}$$

    Notice that the last one is a *complete description*, namely the *outcome* HT.
  - $A \cup B = \{HH, HT, TH\}$, or in words:

    $$\begin{aligned} &\text{"first was heads" OR "exactly one heads"} \\ &= \text{"starts with heads, else it's tails-then-heads"} \end{aligned}$$

Exercise - Coin flipping: counting subsets

Flip a fair coin five times and record the results.

How many elements are in the sample space? (How big is S?)
How many events are there? (How big is $\mathcal{F}$?)

Solution

03 Theory

New events from old

Given two events A and B, we can form new events using set operations:

$$A \cup B \quad \longleftrightarrow \quad \text{“event } A \text{ OR event } B\text{”}$$

$$A \cap B \quad \longleftrightarrow \quad \text{“event } A \text{ AND event } B\text{”}$$

$$A^c \quad \longleftrightarrow \quad \textbf{not} \text{ event } A$$

We also use these terms for events A and B:

- They are **mutually exclusive** when $A \cap B = \emptyset$, that is, they have *no elements in common*.
- They are **collectively exhaustive** $A \cup B = S$, that is, when they jointly *cover all possible outcomes*.

- In probability texts, sometimes $A \cap B$ is written "$A \cdot B$" or even (frequently!) "AB".

Rules for sets

Algebraic rules

- Associativity: $(A \cup B) \cup C = A \cup (B \cup C)$. Analogous to $(A + B) + C = A + (B + C)$.
- Distributivity: $A \cap (B \cup C) = (A \cap B) \ \cup \ (A \cap C)$. Analogous to $A(B + C) = AB + AC$.

De Morgan's Laws

- $(A \cup B)^c = A^c \cap B^c$
- $(A \cap B)^c = A^c \cup B^c$
  In other words: you can distribute " ^{c} " but must simultaneously do a switch $\cap \leftrightarrow \cup$.

Probability models

04 Theory

Axioms of probability

A **probability measure** is a function $P : \mathcal{F} \to \mathbb{R}$ satisfying:

Kolmogorov Axioms:

- **Axiom 1:** $P[A] \geq 0$ for every event A
  (probabilities are not negative!)
- **Axiom 2:** $P[S] = 1$
  (probability of "anything" happening is 1)
- **Axiom 3:** additivity for any *countable collection* of *mutually exclusive* events:

$$P[A_1 \cup A_2 \cup A_3 \cup \cdots] = P[A_1] + P[A_2] + P[A_3] + \cdots$$

$$\text{when:} \quad A_i \cap A_j = \emptyset \quad \text{for all } \ i \neq j$$

- Notation: we write $P[A]$ instead of $P(A)$, even though P is a function, to emphasize the fact that A is a set.

Probability model

A **probability model** or **probability space** consists of a triple $(S, \mathcal{F}, P)$:

- S the sample space
- $\mathcal{F}$ the set of valid events, where every $A \in \mathcal{F}$ satisfies $A \subset S$
- $P : \mathcal{F} \to \mathbb{R}$ a probability measure satisfying the Kolmogorov Axioms

Finitely many exclusive events

It is a consequence of the Kolmogorov Axioms that additivity also works for finite collections of events:

$$P[A \cup B] = P[A] + P[B]$$

$$P[A_1 \cup \cdots \cup A_n] = P[A_1] + \cdots + P[A_n]$$

Inferences from Kolmogorov

A probability measure satisfies these rules.
They can be deduced from the Kolmogorov Axioms.

- **Negation:** Can you find $P[A^c]$ but not $P[A]$? Use negation:

$$P[A] = 1 - P[A^c]$$

- **Monotonicity:** Probabilities grow when outcomes are added:

$$A \subset B \quad \gg\!\gg \quad P[A] \leq P[B]$$

- **Inclusion-Exclusion:** A trick for resolving unions:

$$P[A \cup B] = P[A] + P[B] - P[A \cap B]$$

(even when A and B are *not exclusive!*)

Inclusion-Exclusion

The principle of inclusion-exclusion generalizes to three events:

$$P[A \cup B \cup C] =$$

$$P[A] + P[B] + P[C] - P[A \cap B] - P[A \cap C] - P[B \cap C] + P[A \cap B \cap C]$$

The same pattern works for any number of events!

The pattern goes: "include singles" then "exclude doubles" then "include triples" then ...

Include, exclude, include, exclude, include, ...

05 Illustration

Example - Lucia is Host or Player

The professor chooses three students at random for a game in a class of 40, one to be Host, one to be Player, one to be Judge. What is the probability that Lucia is either Host or Player?

Solution

1. Set up the probability model.
    - Label the students 1 to 40. Write L for Lucia's number.
    - *Outcomes:* assignments such as $(H, P, J) = (2, 5, 8)$
      These are ordered triples with *distinct* entries in $1, 2, \dots, 40$.
    - *Sample space:* S is the collection of all such distinct triples
    - *Events:* any subset of S
    - *Probability measure*: assume all outcomes are equally likely, so $P[(i,j,k)] = P[(r,l,p)]$ for all i, j, k, r, l, p
    - In total there are $40 \cdot 39 \cdot 38$ triples of distinct numbers.
    - Therefore $P[(i,j,k)] = \frac{1}{40 \cdot 39 \cdot 38}$ for any *specific* outcome (i,j,k).
    - Therefore $P[A] = \frac{|A|}{40 \cdot 39 \cdot 38}$ for any event A. (Recall $|A|$ is the *number* of outcomes in A.)
2. Define the desired event.
    - Want to find $P[\text{“Lucia is Host or Player”}]$
    - Define $A =$ “Lucia is Host” and $B =$ “Lucia is Player”. Thus:

$$A = \big\{(L,j,k) \mid \text{any } j,k\big\}, \qquad B = \big\{(i,L,k) \mid \text{any } i,k\big\}$$

    - So we seek $P[A \cup B]$.
3. Compute the desired probability.
    - Importantly, $A \cap B = \emptyset$ (mutually exclusive).
      There are no outcomes in S in which Lucia is *both* Host and Player.
    - By *additivity*, we infer $P[A \cup B] = P[A] + P[B]$.
    - Now compute $P[A]$.
        - There are $39 \cdot 38$ ways to choose j and k from the students besides Lucia.
        - Therefore $|A| = 39 \cdot 38$.
        - Therefore:

$$P[A] \quad \gg\gg \quad \frac{|A|}{40 \cdot 39 \cdot 38} \quad \gg\gg \quad \frac{39 \cdot 38}{40 \cdot 39 \cdot 38} \quad \gg\gg \quad \frac{1}{40}$$

- Now compute $P[B]$. It is similar: $P[B] = \frac{1}{40}$.
- Finally compute that $P[A] + P[B] = \frac{1}{20}$, so the answer is:

$$P[A \cup B] \quad \gg\gg \quad P[A] + P[B] \quad \gg\gg \quad \frac{1}{20}$$

Example - iPhones and iPads

At Mr. Jefferson's University, 25% of students have an iPhone, 30% have an iPad, and 60% have neither.

What is the probability that a randomly chosen student has *some* iProduct? (Q1)

What about *both*? (Q2)

Solution

1. Set up the probability model.
    - A student is chosen at random: an *outcome* is the chosen student.
    - *Sample space* S is the set of all students.
    - Write $O =$ "has iPhone" and $A =$ "has iPad" concerning the chosen student.
    - All students are equally likely to be chosen: therefore $P[E] = \frac{|E|}{|S|}$ for any event E.
    - Therefore $P[O] = 0.25$ and $P[A] = 0.30$.
    - Furthermore, $P[O^c A^c] = 0.60$. This means 60% have "not iPhone AND not iPad".
2. Define the desired event.
    - Q1: desired event $= O \cup A$
    - Q2: desired event $= OA$
3. Compute the probabilities.
    - We do not believe O and A are exclusive.
    - Try: apply inclusion-exclusion:

$$P[O \cup A] = P[O] + P[A] - P[OA]$$

    - We know $P[O] = 0.25$ and $P[A] = 0.30$. So this formula, with given data, RELATES Q1 and Q2.
    - Notice the complements in $O^c A^c$ and try *Negation*.
    - *Negation:*

$$P[(OA)^c] = 1 - P[OA]$$

    DOESN'T HELP.

- Try again: *Negation:*

$$P[(O^cA^c)^c] = 1 - P[O^cA^c]$$

- And De Morgan (or a Venn diagram!):

$$(O^cA^c)^c \quad \gg\gg \quad O \cup A$$

- Therefore:

$$P[O \cup A] \quad \gg\gg \quad P[(O^cA^c)^c]$$

$$\gg\gg \quad 1 - P[O^cA^c] \quad \gg\gg \quad 1 - 0.6 = 0.4$$

- We have found Q1: $P[O \cup A] = 0.40$.
- Applying the RELATION from inclusion-exclusion, we get Q2:

$$P[O \cup A] = P[O] + P[A] - P[OA]$$

$$\gg\gg \quad 0.40 = 0.25 + 0.30 - P[OA]$$

$$\gg\gg \quad P[OA] = 0.15$$

Conditional probability

06 Theory

Conditional probability

The **conditional probability** of "B given A" is defined by:

$$P[B \mid A] = \frac{P[B \cap A]}{P[A]}$$

This conditional probability $P[B \mid A]$ represents the probability of event B taking place *given the assumption* that A took place. (All within the given probability model.)

By letting the actuality of event A be taken as a fixed hypothesis, we can define a *conditional probability measure* by plugging events into the slot of B:

$$P[- \mid A] = \frac{P[- \cap A]}{P[A]}$$

It is possible to verify each of the Kolmogorov axioms for this function, and therefore $P[- \mid A]$ itself defines a bona fide *probability measure*.

Conditioning

What does it really mean?

Conceptually, $P[B \mid A]$ corresponds to *creating a new experiment* in which we run the old experiment and record data *only those times that A happened*. Or, it corresponds to finding ourselves with *knowledge* or *data* that A happened, and we seek our best estimates of the likelihoods of other events, based on our existing model and the actuality of A.

Mathematically, $P[B \mid A]$ corresponds to *restricting* the probability function to outcomes in A, and *renormalizing* the values (dividing by $p[A]$) so that the total probability of all the outcomes (in A) is now 1.

The definition of conditional probability can also be turned around and reinterpreted:

Multiplication rule

$$P[AB] = P[A] \cdot P[B \mid A]$$

"The probability of A AND B equals the probability of A *times* the probability of B -given-A."

This principle generalizes to any events in sequence:

Generalized multiplication rule

$$P[A_1A_2A_3] = P[A_1] \cdot P[A_2 \mid A_1] \cdot P[A_3 \mid A_1A_2]$$

$$P[A_1 \cdots A_n] = P[A_1] \cdot P[A_2 \mid A_1] \cdot P[A_3 \mid A_1A_2] \, \cdots \, P[A_n \mid A_1 \, \cdots \, A_{n-1}]$$

The generalized rule can be verified like this. First substitute A_2 for B and A_1 for A in the original rule. Now repeat, substituting A_3 for B and A_1A_2 for A in the original rule, and combine with the first one, and you find the rule for triples. Repeat again with A_4 and $A_1A_2A_3$, combine with the triples, and you get quadruples.

07 Illustration

Exercise - Simplifying conditionals

Let $A \subset B$. Simplify the following values:

$$P[A \mid B], \quad P[A \mid B^c], \quad P[B \mid A], \quad P[B \mid A^c]$$

Solution

Coin flipping: at least 2 heads

Flip a fair coin 4 times and record the outcomes as sequences, like $HHTH$.

Let $A_{\geq 2}$ be the event that there are at least two heads, and $A_{\geq 1}$ the event that there is at least one heads.

First let's calculate $P[A_{\geq 2}]$.

Define A_2, the event that there were exactly 2 heads, and A_3, the event of exactly 3, and A_4 the event of exactly 4. These events are exclusive, so:

$$P[A_{\geq 2}] = P[A_2 \cup A_3 \cup A_4] \quad \gg\gg \quad P[A_2] + P[A_3] + P[A_4]$$

Each term on the right can be calculated by counting:

$$P[A_2] = \frac{|A_2|}{2^4} \quad \gg\gg \quad \frac{\binom{4}{2}}{16} \quad \gg\gg \quad \frac{6}{16}$$

$$P[A_3] = \frac{|A_3|}{2^4} \quad \gg\gg \quad \frac{\binom{4}{1}}{16} \quad \gg\gg \quad \frac{4}{16}$$

$$P[A_4] = \frac{|A_4|}{2^4} \quad \gg\gg \quad \frac{\binom{4}{0}}{16} \quad \gg\gg \quad \frac{1}{16}$$

Therefore, $P[A_{\geq 2}] = \frac{11}{16}$.

Now suppose we find out that "at least one heads definitely came up". (Meaning that we know $A_{\geq 1}$.) For example, our friend is running the experiment and tells us this fact about the outcome.

Now what is our estimate of likelihood of $A_{\geq 2}$?

The formula for conditioning gives:

$$P[A_{\geq 2} \mid A_{\geq 1}] = \frac{P[A_{\geq 2} \cap A_{\geq 1}]}{P[A_{\geq 1}]}$$

Now $A_{\geq 2} \cap A_{\geq 1} = A_{\geq 2}$. (Any outcome with at least two heads automatically has at least one heads.) We already found that $P[A_{\geq 2}] = \frac{11}{16}$. To compute $P[A_{\geq 1}]$ we simply *add* the probability $P[A_1]$, which is $\frac{4}{16}$, to get $P[A_{\geq 1}] = \frac{15}{16}$.

Therefore:

$$P[A_{\geq 2} \mid A_{\geq 1}] = \frac{11/16}{15/16} \quad \gg\gg \quad \frac{11}{15}$$

Example - Flip a coin, then roll dice

Flip a coin. If the outcome is heads, roll two dice and add the numbers. If the outcome is tails, roll a single die and take that number. What is the probability of getting a tails AND a number at least 3?

Solution

This "two-stage" experiment lends itself to a solution using conditional probability.

1. Label the events of interest.
   - Let H and T be the events that the coin showed heads and tails, respectively.
   - Let $A_1, \ldots, A_{12}$ be the events that the final number is 1, …, 12, respectively.
   - The value we seek is $P[TA_{\geq 3}]$.
2. Observe known (conditional) probabilities.
   - We know that $P[H] = 0.5$ and $P[T] = 0.5$.
   - We know that $P[A_5 \mid T] = \frac{1}{6}$, for example, or that $P[A_1 \mid H] = \frac{1}{12}$.
3. Apply "multiplication" rule.
   - This rule gives:

     $$P[TA_{\geq 3}] = P[T] \cdot P[A_{\geq 3} \mid T]$$

   - We know $P[T] = 0.5$ and can see by counting that $P[A_{\geq 3} \mid T] = 0.5$.
   - Therefore $P[TA_{\geq 3}] = 0.25$.